Expanded held-out test sample

Downloaded and audited **1,600 new images** from Tiny-GenImage. Each new test contains 400 real and 400 generated images, compared with 40 + 40 in the original pilot. All original training, validation and test images were excluded.

New test	Generated sources	Real	Generated	Total
Seen	BigGAN 200 + GLIDE 200	400	400	800
Unseen	Midjourney	400	400	800

File hashes and decoded RGB pixel hashes are unique across the new sample and disjoint from the 480-image pilot. One decision now changes a run's FPR by 0.0025, compared with 0.025 previously. The same six fitted checkpoints and the original threshold 0.5 were used; there was no retraining.

Measured results on the new original images

Model	Test	F1 (SD)	FPR (SD)	AUROC (SD)
ResNet-50	Seen	0.955 (0.018)	0.033 (0.014)	0.991 (0.003)
CLIP linear	Seen	0.884 (0.005)	0.251 (0.012)	0.977 (0.002)
ResNet-50	Unseen	0.157 (0.106)	0.034 (0.023)	0.670 (0.029)
CLIP linear	Unseen	0.640 (0.005)	0.255 (0.009)	0.738 (0.007)

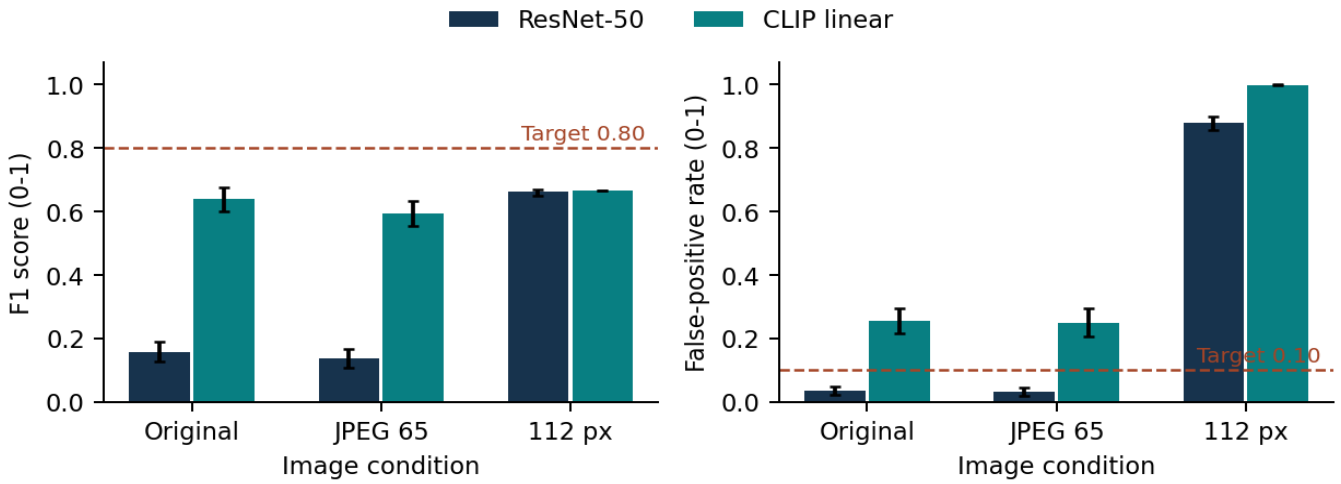

Figure 1. New unseen test, n = 800; means over seeds 42, 43 and 44. Error bars are 95% stratified image-bootstrap intervals conditional on the three checkpoint pairs. JPEG 65 and 112 px are separate processing conditions.

On original unseen images, mean F1 differed by 0.483 (CLIP minus ResNet; 95% CI [0.437, 0.527]; paired permutation p = 0.0001). This later test was collected after inspecting the first pilot. It increases the sample size; native-size, compression and category confounds remain unresolved.

The original pilot sampled the first 600 validation rows; this draw uses the full 7,000-row frame. Neither model reaches unseen F1 0.80, and CLIP's mean FPR of 0.255 exceeds the 0.10 target.

Source: TheKernel01/Tiny-GenImage, Hugging Face; revision 89c4fe9efd0e. The complete 7,000-row validation split was audited before random selection with seed 20261017. The evidence ZIP contains the manifest, predictions, audit, scripts and full per-run statistics. Dataset images retain the upstream CC BY-NC-SA 4.0 terms.